

# CMSC 510 – Fall 2018



---

## Homework Assignment 1

Announced: 9/26

Due: Monday, 10/15, 3pm



# Parametric Modeling

---

- Maximum likelihood (actually, maximum a posteriori) estimation of model parameters, for two-class classification
- Part A:  $p(x|y)$  is Gaussian  
we have two distributions:
  - $p(x|y=\text{classA})$
  - $P(x|y=\text{classB})$
- Part B:  $p(y|x)$  as in Logistic Regression



# Parametric Modeling

---

- We will use PyMC3
  - Statistical library for Python 3
  - We specify a statistical model with some unknown parameters
    - E.g.  $P(x_k|\theta_i) \Leftrightarrow x \sim \text{Normal}(\text{mean}, \text{covariance})$   
mean and covariance are unknown
  - We tie it to actual observed data
    - E.g. samples (feature vectors)  $x_k$
- And then we estimate the parameters
  - E.g. we find some good values for mean and covariance

# Parametric Modeling

- **Part A:**  $p(x|y)$  is Gaussian

$p(x|y=\text{classA})$  and  $P(x|y=\text{classB})$

- See L06:

- Maximum a posteriori (MAP):

- Finding  $\theta_i$  that maximizes  $P(\theta_i|S_i) \sim P(S_i|\theta_i)P(\theta_i)$
- Where  $P(S_i|\theta_i)P(\theta_i) = \prod_k P(x_k|\theta_i)P(\theta_i)$

- In our case  $P(x_k|\theta_i) = \text{Normal}(\text{mean}, \text{covariance})$


- That is,  $\theta_i$  represents the mean and the covariance matrix of the gaussian

- In PyMC3, we model that as (unknown, known):

- `x1_obs = pm.MvNormal('x1', mu=mu1, chol=chol, observed = x1);`

- `x0_obs = pm.MvNormal('x0', mu=mu0, chol=chol, observed = x0);`

- Note: different means, same covariance (given here as Cholesky Decomposition of Cov. Mx)



# Parametric Modeling


- **Part B:**  $p(y|x)$  as in Logistic Regression

- $p(x|y=\text{classA})$  and  $P(x|y=\text{classB})$

- See L06:

- We have:

$$P(y_i | x_i, w) = a(y_i w^T x_i) = \frac{1}{1 + e^{-y_i w^T x_i}}$$



- We will also use a prior  $P(w)$  (a Gaussian prior)
- $Y$  is modelled as a Bernoulli distribution

In probability theory and statistics, the **Bernoulli distribution**, named after Swiss mathematician Jacob Bernoulli,<sup>[1]</sup> is the discrete probability distribution of a random variable which takes the value 1 with probability  $p$  and the value 0 with probability  $q = 1 - p$ , that is, the probability distribution of any single experiment that asks a yes–no question; the question results in a boolean-valued outcome, a single bit of information whose value is success/yes/true/one with probability  $p$  and failure/no/false/zero with probability  $q$ . It can be used to represent a coin toss where 1 and 0 would represent "head" and "tail" (or vice versa), respectively. In particular, unfair coins would have  $p \neq 0.5$ .

- In PyMC3, that's:

- `Y_obs=pm.Bernoulli('Y_obs', p=prob, observed = Y);`

# PyMC3

- We will use PyMC3
  - There are two .py files on BB, with a working program for 2-dimensional (2 features) fake dataset
  - That should be your starting point
- Your task is to perform experiments with the two approaches on MNIST dataset
  - `from keras.datasets import mnist`
  - `(x_train, y_train), (x_test, y_test) = mnist.load_data()`
- This is 10-class classification problem:  
hand-written digits
- It has 5,000 samples per class
  - During development, you can randomly select a subset of samples, to speed up the calculations





# MNIST problem

---

- MNIST: a 10-class classification problem
- Convert it into 2-class problem by: taking last digit of your V# (class A), taking last different digit of your V# (class B)
  - E.g. V# V00078965: 6-vs-5, V00078966: 9-vs-6
- Train both methods on the training set, then apply them to test set, and calculate the % of error predictions
  - For method A, estimate means and covariance  $\Sigma$  of the gaussians from training data, then use the means and covariance  $\Sigma$  to classify samples from the test set
    - how? use formula for  $p(x|y=A)$ , and  $p(x|y=B)$ , and pick class with higher probability
  - For method B, estimate  $w, b$  from training data, then use  $w, b$  to classify samples from the test set
- Try the same for smaller training set (e.g. using only random 50%, 10%, 5% of training samples), but always use the same full test set
  - see if/how much the error rate on the test set increases



# Returning the Assignment

---

- Solution code should be written by you and you only (no web/book/friend/etc. code)
  - You can freely use the code provided on BB as your starting point
- Upload through Blackboard
  - A report in PDF
    - Results of tests of the two methods on MNIST dataset, for decreasing training set sizes (include you V#, and what are your two digits defining the two-class problem).
  - Code in python for solving the MNIST classification problem (for full size of the training set):
    - `FamilyNameFirstName-PartA.py`
    - `FamilyNameFirstName-PartB.py`
      - The files should have your name in a comment at the top